

Data Science Lab

Vasudevan T V

Matrix Operations

- ▶ Using Vectorisation
- ▶ Here various matrix operations are performed **without using loops**
- ▶ For this, we can use various functions in the built in package **numpy**

```
# Matrix Addition
```

```
>>> import numpy
```

```
>>> matrix1=numpy.matrix([[1,2],[3,4]])
```

```
>>> matrix2=numpy.matrix([[4,3],[2,1]])
```

```
>>> matrix3=numpy.add(matrix1,matrix2)
```

```
>>> print(matrix3)
```

```
[[5 5]
```

```
 [5 5]]
```

Matrix Operations

Matrix Subtraction

```
>>> import numpy
>>> matrix1=numpy.matrix([[2,2],[2,2]])
>>> matrix2=numpy.matrix([[1,1],[1,1]])
>>> matrix3=numpy.subtract(matrix1,matrix2)
>>> print(matrix3)
[[1 1]
 [1 1]]
```

Matrix Multiplication

```
>>> import numpy
>>> matrix1=numpy.matrix([[2,2],[2,2]])
>>> matrix2=numpy.matrix([[1,1],[1,1]])
>>> matrix3=numpy.matmul(matrix1,matrix2)
>>> print(matrix3)
[[4 4]
 [4 4]]
```

Matrix Operations

Scalar Multiplication

```
>>> import numpy
>>> matrix1=numpy.matrix([[2,2],[2,2]])
>>> matrix2=2*matrix1
>>> print(matrix2)
[[4 4]
 [4 4]]
```

Matrix Transpose

```
>>> import numpy
>>> matrix1=numpy.matrix([[1,2],[3,4]])
>>> print(matrix1)
[[1 2]
 [3 4]]
>>> matrix2=numpy.transpose(matrix1)
>>> print(matrix2)
[[1 3]
 [2 4]]
```

Matrix Transformations

- ▶ We can use matrices for performing various geometric transformations such as **translation**, **rotation**, **scaling** etc.
- ▶ **Translation** is the process of moving an object to a different position
- ▶ **Rotation** is the process of changing the angle of the object
- ▶ **Scaling** is the process of changing the size of objects

Matrix Transformations

► Translation Matrix

$$\begin{bmatrix} 1 & 0 & T_x \\ 0 & 1 & T_y \\ 0 & 0 & 1 \end{bmatrix}$$

► Program

```
import numpy
def translationMatrix(tx=0, ty=0):
    return numpy.matrix([[1,0,tx],
                          [0,1,ty],
                          [0,0, 1]])
matrix=translationMatrix(1,1)
print(matrix)
```

Matrix Transformations

► Rotation Matrix

$$\begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

► Program

```
import numpy
def rotationMatrix(degree):
    theta = numpy.radians(degree)
    c,s=numpy.cos(theta),numpy.sin(theta)
    return numpy.matrix([[c, -s, 0],
                          [s,  c, 0],
                          [0,  0, 1]])

matrix=rotationMatrix(30)
print(matrix)
```

Matrix Transformations

► Scaling Matrix

$$\begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

► Program

```
import numpy
def scalingMatrix(sx=0, sy=0):
    return numpy.matrix([[sx,0,0],
                          [0,sy,0],
                          [0, 0,1]])
matrix=scalingMatrix(2,2)
print(matrix)
```


Singular Value Decomposition (SVD)

- ▶ It is the process of decomposing a matrix into 3 components which are also matrices
- ▶ A matrix M is decomposed into 3 matrices U , S and V
- ▶ If M is a real matrix, U and V are orthogonal matrices and S is a diagonal matrix
- ▶ The advantage of such a decomposition is that we can do the subsequent matrix operations faster
- ▶ Applications - solving homogeneous linear equations, pattern recognition, natural language processing, weather prediction, machine learning etc.

Singular Value Decomposition (SVD)

► Program

```
# Imports matrix, matmul and diag functions only
from numpy import matrix
from numpy import matmul
from numpy import diag
# Imports svd fn from linalg(linear algebra) submodule of
# scipy module
from scipy.linalg import svd
# define a matrix
A = matrix([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
print(A)
# Singular-value decomposition
# A is decomposed into 3 matrices U, a diagonal matrix
# and V
# Here S contains only the diagonal elements of the
# diagonal matrix
U, S, V = svd(A)
```

Singular Value Decomposition (SVD)

► Program - continued

```
print(U)
print(S)
print(V)
# create diagonal matrix from diagonal elements
Sigma = diag(S)
print(Sigma)
# reconstruct matrix
B = matmul(U,matmul(Sigma,V))
print(B)
```

Singular Value Decomposition (SVD)

► Output

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
[[-0.21483724  0.88723069  0.40824829]
 [-0.52058739  0.24964395 -0.81649658]
 [-0.82633754 -0.38794278  0.40824829]]
[1.68481034e+01 1.06836951e+00 4.41842475e-16]
[[-0.47967118 -0.57236779 -0.66506441]
 [-0.77669099 -0.07568647  0.62531805]
 [-0.40824829  0.81649658 -0.40824829]]
[[1.68481034e+01 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 1.06836951e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 4.41842475e-16]]
[[1. 2. 3.]
 [4. 5. 6.]
 [7. 8. 9.]]
```

Histogram

- ▶ Write a python program to plot a histogram of marks obtained by students in a class
- ▶ Marks - 22,87,5,43,56,73,55,54,11,20,51,5,79,31,27

```
# imports pyplot, a module used in the package matplotlib
# to plot various figures
from matplotlib import pyplot
# imports array() from numpy package
from numpy import array
# subplots() specify the number of plots in the figure
# first argument is number of rows
# second argument is number of columns
# This function returns a tuple containing figure and axes
# objects
# These objects are assigned to fig and ax
# They are needed for changing figure level and axes level
# attributes
fig,ax = pyplot.subplots(1,1)
```

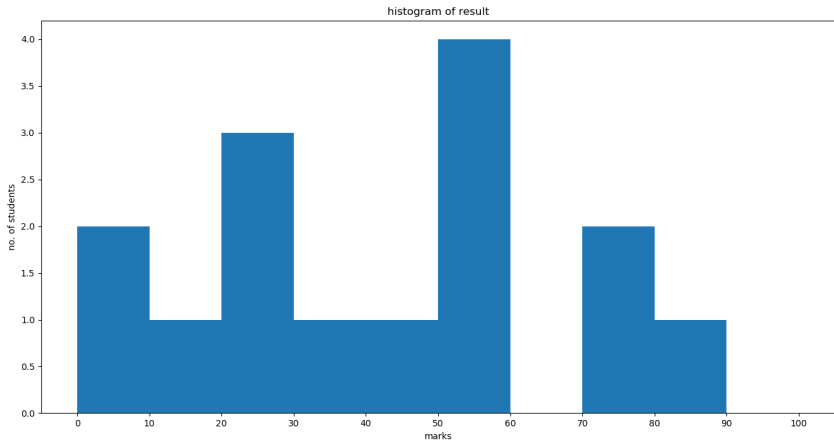
Histogram

► Program - continued

```
a = array([22,87,5,43,56,73,55,54,11,20,51,5,79,31,27])
# Draws a histogram, first argument is the array of
# numbers, second argument bins are intervals of values
ax.hist(a,bins=[0, 10, 20, 30, 40, 50, 60, 70, 80,90,100])
ax.set_title("histogram of result")
ax.set_xticks([0, 10, 20, 30, 40, 50, 60, 70, 80, 90,100])
ax.set_xlabel('marks')
ax.set_ylabel('no. of students')
# Shows the plot
pyplot.show()
```

Histogram

► Output



Histogram

- ▶ Write a python program to draw a histogram of petal length in the iris data set
- ▶ Program

```
from matplotlib import pyplot
# imports pandas package, used for data analysis
import pandas
# reads the csv file into a data frame
# A data frame is a table with rows and columns
df = pandas.read_csv('iris.csv')
fig,ax = pyplot.subplots(1,1)
```



Histogram

- ▶ Write a python program to draw a histogram of petal length in the iris data set
- ▶ Program - continued

```
# plots the histogram of petal length attribute
# By default bins = 10
df['petal.length'].plot(kind='hist', edgecolor="black",
bins=49)
ax.set_title("Histogram")
ax.set_xticks([1.0,1.5,2.0,2.5,3.0,3.5,4.0,4.5,5.0,5.5,
6.0,6.5,7.0])
ax.set_yticks([0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,
17,18,19])
ax.set_xlabel('Petal Length')
pyplot.show()
```

Histogram

► Output




Scatter Plot

- ▶ Write a python program to draw a scatterplot that shows the relationship between rollnos and marks of students in a class
- ▶ rollnos = [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]
- ▶ marks = [22,87,5,43,56,73,55,54,11,20,51,5,79,31,27]

```
from matplotlib import pyplot
rollnos = [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]
marks = [22,87,5,43,56,73,55,54,11,20,51,5,79,31,27]
fig,ax = pyplot.subplots(1,1)
# Draws a scatterplot, first argument is x axis values,
# second argument is y axis values
ax.scatter(rollnos, marks)
ax.set_title("Scatter Plot")
ax.set_xticks([1,2,3,4,5,6,7,8,9,10,11,12,13,14,15])
ax.set_yticks([25,50,75,100])
ax.set_xlabel('rollnos')
ax.set_ylabel('marks')
pyplot.show()
```

Scatter Plot

► Output



Scatter Plot

- ▶ Write a python program to draw a scatterplot that shows the relationship between petal length and petal width in the iris data set
- ▶ Program

```
from matplotlib import pyplot
import pandas
df = pandas.read_csv('iris.csv')
fig, ax = pyplot.subplots(1,1)
# Creates a dictionary of colour values of each species
colors = {'Setosa':'red', 'Versicolor':'green',
'Virginica':'blue'}
```


Scatter Plot

- ▶ Write a python program to draw a scatterplot that shows the relationship between petal length and petal width in the iris data set
- ▶ Program - continued

```
# Groups the data based on species values
grouped = df.groupby('species')
# group represents the grouped data frame
# draws the scatter plot for each group
for key, group in grouped:
    group.plot(ax=ax, kind='scatter', x='petal.length',
               y='petal.width', label=key, color=colors[key])
ax.set_title("Scatter Plot")
ax.set_xlabel('Petal Length')
ax.set_ylabel('Petal Width')
pyplot.show()
```

Scatter Plot

► Output



Classification Using kNN Algorithm

- ▶ Given a data set of 15 food items (food.csv) having 4 features - ingredient, sweetness, crunchiness and food type. Write a R program to predict the food type of tomato using kNN algorithm.

```
$ R
```

```
R version 3.3.3 (2017-03-06) -- "Another Canoe"
```

```
.....
```

```
# Read the csv file into a data frame
```

```
> food=read.csv("food.csv")
```


Classification Using kNN Algorithm

```
# Prints food data frame
```

```
> food
```

	Ingredient	Sweetness	Crunchiness	FoodType
1	apple	10	9	fruit
2	bacon	1	4	protein
3	banana	10	1	fruit
4	carrot	7	10	vegetable
5	celery	3	10	vegetable
6	cheese	1	1	protein
7	cucumber	2	8	vegetable
8	fish	3	1	protein
9	grape	8	5	fruit
10	green bean	3	7	vegetable
11	lettuce	1	9	vegetable
12	nuts	3	6	protein
13	orange	7	3	fruit
14	pear	10	7	fruit
15	shrimp	2	3	protein

Classification Using kNN Algorithm

```
# Creates a data frame of food item tomato
> tomato=data.frame(ingredient="tomato",sweetness=6,
crunchiness=4)
# prints the tomato data frame
> tomato
```

	Ingredient	Sweetness	Crunchiness
1	tomato	6	4

```
# Create a data frame of second and third columns of food
> food1=food[,2:3]
```

Classification Using kNN Algorithm

```
> food1
```

	Sweetness	Crunchiness
1	10	9
2	1	4
3	10	1
4	7	10
5	3	10
6	1	1
7	2	8
8	3	1
9	8	5
10	3	7
11	1	9
12	3	6
13	7	3
14	10	7
15	2	3

Classification Using kNN Algorithm

```
# Create a data frame of second and third columns of
# tomato
> tomato1=tomato[,2:3]
> tomato1
  Sweetness Crunchiness
1          6           4
# Load package class which contains knn()
> library(class)
# Use knn() and store the prediction in pred
# argument 1 is the data frame containing training data
# argument 2 is the data frame containing test data
# argument 3 is a vector that show the class of each item
# in the training data, argument 4 is the value of k
> pred=knn(food1,tomato1,food$FoodType,k=1)
> pred
[1] fruit
Levels: fruit protein vegetable
```

Classification Using kNN Algorithm

- ▶ Diagnosing Breast Cancer With The kNN Algorithm
- ▶ The data includes 569 examples of cancer biopsies, each with 32 features
- ▶ One feature is an identification number, another is the cancer diagnosis, and 30 are numeric-valued laboratory measurements
- ▶ The diagnosis is coded as "M" to indicate malignant or "B" to indicate benign
- ▶ The other 30 numeric measurements comprise the mean, standard error, and worst(that is, largest) value for 10 different characteristics of the digitized cell nuclei
- ▶ These include Radius, Texture, Perimeter, Area etc.

Classification Using kNN Algorithm

► Diagnosing Breast Cancer With The kNN Algorithm

```
$ R
R version 3.3.3 (2017-03-06) -- "Another Canoe"
.....
# Loads class package containing knn()
> library(class)
# Loads gmodels package containing CrossTable()
> library(gmodels)
# Read the csv file into a data frame
> wbcd = read.csv("wisc_bc_data.csv")
# Define normalize fn for performing min max normalisation
# This will transform the values of all features to a
# range between 0 and 1
> normalize <- function(x)
{
  return ((x - min(x)) / (max(x) - min(x)))
}
```

Classification Using kNN Algorithm

► Diagnosing Breast Cancer With The kNN Algorithm

```
# Apply this function to our data frame
> wbcd_n = as.data.frame(lapply(wbcd[3:31], normalize))
# Training Data
> wbcd_train = wbcd_n[1:469, ]
# Test data
> wbcd_test = wbcd_n[470:569, ]
# Training Labels
> wbcd_train_labels = wbcd[1:469, 2]
# Test Labels
> wbcd_test_labels = wbcd[470:569, 2]
```

Classification Using kNN Algorithm

► Diagnosing Breast Cancer With The kNN Algorithm

```
# Prediction of Breast Cancer using knn()
> wbcd_test_pred = knn(wbcd_train, wbcd_test,
wbcd_train_labels,21)
# Print the prediction results
> wbcd_test_pred
 [1] B B B B B B B B B B B M B B B B B B B M B B B B M B B..
[38] B B B M B B M B B B M M B B B M B B B B B B B B B B..
[75] B B B B B B B B B B B B B B B B B B B M M M M M M B
Levels: B M
```


Classification Using kNN Algorithm

► Diagnosing Breast Cancer With The kNN Algorithm

```
# Analysis of Prediction
```

```
> CrossTable(x = wbcd_test_labels, y = wbcd_test_pred,  
prop.chisq=FALSE)
```

```
Cell Contents
```

```
|-----|  
|                                     N |  
|           N / Row Total |  
|           N / Col Total |  
|           N / Table Total |  
|-----|
```

Classification Using kNN Algorithm

► Diagnosing Breast Cancer With The kNN Algorithm

Total Observations in Table: 100

	wbcd_test_pred		
wbcd_test_labels	B	M	Row Total
----- ----- ----- -----			
B	77	0	77
	1.000	0.000	0.770
	0.975	0.000	
	0.770	0.000	
----- ----- ----- -----			
M	2	21	23
	0.087	0.913	0.230
	0.025	1.000	
	0.020	0.210	
----- ----- ----- -----			
Column Total	79	21	100
	0.790	0.210	
----- ----- ----- -----			

Classification Using Naive Bayes Algorithm

- Write a R program to predict the species of iris data set using Naive Bayes algorithm and evaluate its performance

```
# Loads e1071 package containing naiveBayes
library(e1071)
# Loads caTools package containing sample.split()
library(caTools)
# Loads gmodels packge containing CrossTable()
library(gmodels)

# Read the csv file into a data frame
iris = read.csv("iris.csv")

# Splitting data into train
# and test data
# set.seed() is used for generating the same sample
# in every execution
# We specify a seed number
set.seed(100)
```

Classification Using Naive Bayes Algorithm

► Program

```
split <- sample.split(iris$species, SplitRatio = 0.7)
iris1 <- subset(iris, split == "TRUE")
iris2 <- subset(iris, split == "FALSE")

iris_train = iris1[,1:4]
iris_test = iris2[,1:4]

iris_train_labels = iris1[,5]
iris_test_labels = iris2[,5]

classifier_cl <- naiveBayes(iris_train,iris_train_labels )
classifier_cl

# Predicting on test data'
iris_test_pred <- predict(classifier_cl, iris_test)
iris_test_pred
```

Classification Using Naive Bayes Algorithm

► Program

```
# Analysis of Prediction
# prop.chisq=FALSE will remove unnecessary chi square
# values
CrossTable(iris_test_labels, iris_test_pred,
prop.chisq=FALSE)
```

► Output

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = iris_train, y = iris_train_labels)
```

A-priori probabilities:

```
iris_train_labels
```

Setosa	Versicolor	Virginica
0.3333333	0.3333333	0.3333333

Classification Using Naive Bayes Algorithm

► Output

Conditional probabilities:

sepal.length

iris_train_labels	[,1]	[,2]
Setosa	5.025714	0.3266072
Versicolor	5.894286	0.5455396
Virginica	6.625714	0.5907836

sepal.width

iris_train_labels	[,1]	[,2]
Setosa	3.445714	0.3567359
Versicolor	2.782857	0.3468223
Virginica	2.985714	0.2658426

petal.length

iris_train_labels	[,1]	[,2]
Setosa	1.471429	0.1808012
Versicolor	4.191429	0.4859021
Virginica	5.608571	0.5083835

Classification Using Naive Bayes Algorithm

► Output

```
                petal.width
iris_train_labels  [,1]      [,2]
      Setosa      0.2285714 0.08934872
      Versicolor 1.3228571 0.20448747
      Virginica   2.0485714 0.28218833
```

```
# For numerical values, conditional probabilities display
# their mean and standard deviation
```

```
[1] Setosa      Setosa      Setosa      Setosa      Setosa      Setosa
[7] Setosa      Setosa      Setosa      Setosa      Setosa      Setosa
[13] Setosa      Setosa      Setosa      Versicolor  Versicolor  Versicolor
[19] Versicolor  Versicolor  Versicolor  Virginica   Versicolor  Versicolor
[25] Versicolor  Versicolor  Versicolor  Versicolor  Versicolor  Versicolor
[31] Versicolor  Virginica   Virginica   Virginica   Virginica   Virginica
[37] Virginica   Versicolor  Virginica   Virginica   Virginica   Virginica
[43] Virginica   Virginica   Virginica
Levels: Setosa Versicolor Virginica
```

Classification Using Naive Bayes Algorithm

► Output

Cell Contents

	N
N / Row Total	
N / Col Total	
N / Table Total	

Total Observations in Table: 45

iris_test_labels	iris_test_pred			Row Total
	Setosa	Versicolor	Virginica	
Setosa	15	0	0	15
	1.000	0.000	0.000	0.333
	1.000	0.000	0.000	
	0.333	0.000	0.000	
Versicolor	0	14	1	15
	0.000	0.933	0.067	0.333
	0.000	0.875	0.071	
	0.000	0.311	0.022	
Virginica	0	2	13	15
	0.000	0.133	0.867	0.333
	0.000	0.125	0.929	
	0.000	0.044	0.289	
Column Total	15	16	14	45
	0.333	0.356	0.311	

Classification Using C5.0 Decision Tree Algorithm

- Write a R program to identify risky bank loans using C5.0 Decision Tree Algorithm and evaluate its performance

```
1 # Use C5.0 Decision Tree algorithm to identify risky bank loans
2 # Also evaluate the performance of the algorithm
3 # Given credit.csv data set containing 1000 bank loan records
4 # Loads C50 package containing C5.0()
5 library(C50)
6 # Loads gmodels package containing CrossTable()
7 library(gmodels)
8 # Read the csv file into a data frame
9 credit <- read.csv("credit.csv")
10 # Training Data, 17th column default is omitted
11 credit_train <- credit[1:900,-17 ]
12 # Test Data, 17th column default is omitted
13 credit_test <- credit[901:1000,-17 ]
14 # Training Labels, containing values of 17th column default
15 credit_train_labels = credit[1:900, 17]
16 # Test Labels, containing values of 17th column default
17 credit_test_labels = credit[901:1000, 17]
```

Classification Using C5.0 Decision Tree Algorithm

► Program

```
18 # C5.0() returns a C5.0 model object and stores it in credit_model
19 # credit_train is a data frame containing training data
20 # credit_train_labels is converted into a factor containing categorical values
21 credit_model <- C5.0(credit_train, as.factor(credit_train_labels))
22 # Prints basic data about the decision tree
23 credit_model
24 # Shows the decision tree and some other information
25 summary(credit_model)
26 # Predicting on test data
27 credit_pred <- predict(credit_model, credit_test)
28 credit_pred
29 # Analysis of Prediction
30 # prop.chisq=FALSE will remove unnecessary chi square values
31 CrossTable(credit_test_labels, credit_pred, prop.chisq=FALSE )
```

Classification Using C5.0 Decision Tree Algorithm

► Output

```
Call:
C5.0.default(x = credit_train, y = as.factor(credit_train_labels))

Classification Tree
Number of samples: 900
Number of predictors: 16

Tree size: 63

Non-standard options: attempt to group attributes

Call:
C5.0.default(x = credit_train, y = as.factor(credit_train_labels))

C5.0 [Release 2.07 GPL Edition]          Sun Jan 30 12:54:58 2022
-----

Class specified by attribute `outcome'

Read 900 cases (17 attributes) from undefined.data

Decision tree:

checking_balance in {unknown,> 200 DM}: no (414/53)
checking_balance in {< 0 DM,1 - 200 DM}:
:...months_loan_duration <= 11:
    :...credit_history in {critical,good,poor,perfect}: no (71/11)
    :    credit_history = very good: yes (6/1)
```

Classification Using C5.0 Decision Tree Algorithm

► Output

Time: 0.0 secs

```
[1] yes no no no no no yes no no no no no no no yes yes no no
[19] no yes no no no no no no yes yes yes no yes no no no no no yes
[37] no no yes no no no no no no no no no no no no yes no no no
[55] yes no no no no yes no no no no no no no no no no no yes no
[73] no yes no no no no no no yes no no yes no no no no yes yes
[91] no no yes yes no no yes no yes yes
```

Levels: no yes

Cell Contents

```
|-----|
|              N |
|      N / Row Total |
|      N / Col Total |
|      N / Table Total |
|-----|
```

Total Observations in Table: 100

credit_test_labels	credit_pred		Row Total
	no	yes	
no	55	13	68
	0.809	0.191	0.680
	0.733	0.520	
	0.550	0.130	
yes	20	12	32
	0.625	0.375	0.320
	0.267	0.480	
	0.200	0.120	
Column Total	75	25	100
	0.750	0.250	

Prediction Using Multiple Linear Regression

- Write a R program to predict medical expenses using multiple linear regression technique and evaluate its performance

```
1 # Predict Medical Expenses using Multiple Linear Regression Technique
2 # Also evaluate its performance
3 # Given insurance.csv data set containing 1338 data items
4 # Our model's dependent variable is expenses , which measures the medical costs
5 # each person charged to the insurance plan for the year
6 # Read the csv file into a data frame
7 insurance <- read.csv("insurance.csv")
8 # Training Data
9 insurance_train <- insurance[1:1000,]
10 #Test Data
11 insurance_test <- insurance[1001:1338,]
12 # lm() returns a multiple linear regression model object
13 # the dependent variable expenses goes to the left of the tilde
14 # the independent variables go to the right, separated by + sign
15 # data specifies the data frame in which these variables can be found
16 # lm() is contained in stats package, which is loaded by default
17 insurance_model <- lm(expenses ~ age + sex + bmi + children + smoker + region, data = insurance_train)
18 # Prints estimated regression coefficients
19 insurance_model
20 # Evaluate Model Performance
21 summary(insurance_model)
22 # Predicting on test data
23 insurance_pred <- predict(insurance_model, insurance_test)
24 insurance_pred
```

Prediction Using Multiple Linear Regression

► Output

```
Call:
lm(formula = expenses ~ age + sex + bmi + children + smoker +
    region, data = insurance_train)

Coefficients:
    (Intercept)          age          sexmale          bmi
      -12083.3           264.3          -288.5           339.9
      children      smokeryes  regionnorthwest  regionsoutheast
         410.2       23832.4          -439.9       -1291.3
  regionsouthwest
     -1263.1

Call:
lm(formula = expenses ~ age + sex + bmi + children + smoker +
    region, data = insurance_train)

Residuals:
    Min       10   Median       30      Max
-11070.0  -2783.7   -926.3   1255.7   25270.9

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -12083.26   1135.37  -10.643 < 2e-16 ***
age           264.26     13.40   19.721 < 2e-16 ***
sexmale      -288.53     377.41   -0.765  0.44474
bmi           339.91     32.57   10.436 < 2e-16 ***
children      410.24     156.89    2.615  0.00906 **
smokeryes    23832.38    475.70   50.099 < 2e-16 ***
regionnorthwest -439.90    543.63   -0.809  0.41861
regionsoutheast -1291.29    534.51   -2.416  0.01588 *
regionsouthwest -1263.15    537.30   -2.351  0.01892 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5934 on 991 degrees of freedom
Multiple R-squared:  0.7569,    Adjusted R-squared:  0.7549
F-statistic: 385.7 on 8 and 991 DF,  p-value: < 2.2e-16
```

Prediction Using Multiple Linear Regression

► Output

1001	1002	1003	1004	1005	1006	1007
27583.4127	27654.6512	1476.9030	9110.7582	6981.4349	6457.5775	5793.6412
1008	1009	1010	1011	1012	1013	1014
34262.0651	3547.7564	10944.6876	7087.9132	29195.2402	15715.3142	11261.9993
1015	1016	1017	1018	1019	1020	1021
6076.8905	11433.6305	1271.4416	5969.6590	15151.3352	4954.9571	12418.8576
1022	1023	1024	1025	1026	1027	1028
28046.1886	35263.5807	-569.5267	14860.4921	3963.8578	25299.6560	-372.3569
1029	1030	1031	1032	1033	1034	1035
11376.1261	4391.7943	31915.8844	36956.8261	5338.1408	23547.3657	16353.7672
1036	1037	1038	1039	1040	1041	1042
9972.1544	29403.8658	33976.6012	3258.4527	2297.3010	30084.2063	231.6483
1043	1044	1045	1046	1047	1048	1049
27175.8495	2822.4706	14552.7358	31888.5500	7804.7607	34265.6128	2146.5830
1050	1051	1052	1053	1054	1055	1056
33649.2764	12075.7742	13517.7324	11126.5799	33977.5689	1909.6591	11119.2972

Prediction Using Multiple Linear Regression

- ▶ **Output** - insurance_model
- ▶ The **Intercept** is the predicted value of **expenses** when the independent variables are equal to zero
- ▶ The other **Coefficients** indicate the estimated increase in **expenses** for an increase of one in each of the features, assuming all other values are held constant
- ▶ For each additional year of **age**, **medical expenses** will be increased by 264.3, when all other features remain constant
- ▶ For each additional **child**, **medical expenses** will be increased by 410.2, when all other features remain constant

Prediction Using Multiple Linear Regression

- ▶ **Output** - `summary(insurance_model)`
- ▶ The **Residuals** section provides summary statistics for the errors in our prediction
- ▶ The **Coefficients** section provides statistics for the errors associated with regression coefficients
- ▶ The **multiple R-squared value** indicates the variation in the **dependent variable**, which is nearly 75 percent

k means Clustering Algorithm

► Program

```
# Write a program to partition the iris data set(given) into different clusters using k-means clustering algorithm.
# Choose k as 3.Check clustering result against species class label.

# set.seed() is used for generating the same sample in every execution
# We specify a seed number
set.seed(100)

# Reads the iris data set into the iris data frame
iris <- read.csv("iris.csv")

# Make a copy of iris data
iris2 <- iris

# Remove the species class label
iris2$species <- NULL

# Clustering with kmeans is performed by kmeans()
# kmeans() is contained within stats package which is loaded by default
# The kmeans() function requires a data frame containing only numeric data and a parameter specifying the desired number of clusters
# This function will return a cluster object that stores cluster information
# The cluster information includes cluster sizes, cluster means, vector of cluster assignments etc.
iris_clusters <- kmeans(iris2, 3)
print(iris_clusters)

# Check clustering result against species class label
# iris_clusters$cluster is a vector of cluster assignments from the kmeans()
# table() performs a tabulation of categorical variable and gives its frequency as output
table(iris$species, iris_clusters$cluster)
```

k means Clustering Algorithm

► Output

```
MES13s-Mac-mini:Data Science Lab mes13$ Rscript 33KMC.R
K-means clustering with 3 clusters of sizes 33, 21, 96
```

Cluster means:

	sepal.length	sepal.width	petal.length	petal.width
1	5.175758	3.624242	1.472727	0.2727273
2	4.738095	2.904762	1.790476	0.3523810
3	6.314583	2.895833	4.973958	1.7031250

Clustering vector:

```
[1] 1 2 2 2 1 1 1 1 2 2 1 1 2 2 1 1 1 1 1 1 1 1 2 2 1 1 1 2 2 1 1 1 2 1 1
[38] 1 2 1 1 2 2 1 1 2 1 2 1 1 3 3 3 3 3 3 3 2 3 3 2 3 3 3 3 3 3 3 3 3 3 3
[75] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 2 3 3 3 3 2 3 3 3 3 3 3 3 3 3 3 3
[112] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
[149] 3 3
```

Within cluster sum of squares by cluster:

```
[1] 6.432121 17.669524 118.651875
(between_SS / total_SS = 79.0 %)
```

Available components:

[1] "cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6] "betweenss"	"size"	"iter"	"ifault"	
	1 2 3			
Setosa	33 17 0			
Versicolor	0 4 46			
Virginica	0 0 50			